



POLITECNICO
MILANO 1863

SCUOLA DI INGEGNERIA INDUSTRIALE
E DELL'INFORMAZIONE

Procedural Decoration of DOOM Levels with Generative Adversarial Networks

TESI DI LAUREA MAGISTRALE IN
INGEGNERIA INFORMATICA

Author: **Akash Aloysius James**

Student ID: 10687690
Advisor: Daniele Loiacono
Co-advisors: Pierluca Lanzi
Academic Year: 2022-23

Abstract

This study examines how to automate the process of populating procedurally generated levels using Conditional Adversarial Networks. The goal is to spawn game objects in a manner that functionally mimics the characteristics of those present in human designed levels without the need to explicitly define an exhaustive set of rules to govern the system. It is meant to enhance the level building process by upgrading the process of including content within the generated DOOM levels. The initial objective is to produce a coherent topology, accompanied by a suitable availability and positioning of game objects with respect to the level layout. This requires addressing complications that impede consistent representation and generation from the previous design. Then the modifications are compiled together to fashion a generative system that can construct the minimum sufficient subset of features used to design a unique level. Finally, a comparison is made between the modified system and the prior architecture to highlight the improvements it provides in regard to what was previously feasible. These samples are evaluated in a manner that contrasts the differences present among them and those available in the repository. The outcome of these experiments should prove a degree of competence in the system's ability for level design and act as a practical alternative for the implementation of procedural content generation for games to come.

Abstract in lingua italiana

Questa tesi esamina come automatizzare il popolamento di livelli generati proceduralmente utilizzando delle Conditional Adversarial Networks. L'obiettivo è generare oggetti di gioco in modo tale da imitare dal punto di vista funzionale le caratteristiche di quelli presenti nei livelli progettati dall'uomo, senza la necessità di definire esplicitamente un insieme esaustivo di regole a questo scopo. L'obiettivo è migliorare il processo di creazione dei livelli, potenziando il processo di inserimento di contenuti all'interno dei livelli generati di DOOM. L'obiettivo iniziale è produrre una topologia coerente, accompagnata da una disponibilità e posizionamento adeguati degli oggetti di gioco rispetto alla struttura del livello. Ciò richiede di affrontare complicazioni che ostacolano la rappresentazione e la generazione di contenuti coerenti a partire dal design originale. Successivamente, le componenti progettate vengono combinate per creare un sistema generativo in grado di costruire il sottoinsieme minimo sufficiente di caratteristiche utilizzate per progettare un livello unico. Infine, viene effettuato un confronto tra il sistema modificato e i precedenti lavori per evidenziare i miglioramenti che offre rispetto a quanto era precedentemente possibile. I livelli generati sono stati valutati in modo da evidenziare le differenze con i livelli disponibili nel repository. Il risultato di questa analisi sperimentale consente di dimostrare un certo grado di competenza del sistema nella progettazione dei livelli, offrendo un'alternativa pratica per l'implementazione della generazione procedurale di contenuti per i giochi futuri.

Contents

Abstract	i
Abstract in lingua italiana	iii
Contents	v
1 Introduction	1
1.1 Background	2
1.2 Problem	2
1.3 Goals	3
1.4 Thesis Structure	3
2 Theory and Motivation	5
2.1 Related Works	5
2.2 State of the Art	6
2.3 Literature Review on Procedural Generation of Game Levels	7
2.4 Theoretical Background	8
2.4.1 Generative Adversarial Networks	8
2.4.2 Deep Convolution Generative Adversarial Networks	9
2.4.3 Wasserstien Generative Adversarial Networks with Gradient Penalty	10
2.4.4 U-Net	11
2.4.5 Conditional Generative Adversarial Networks	11
3 Representation and Prior Design	13
3.1 DOOM	13
3.2 WAD	14
3.2.1 Format	14
3.2.2 Structure	14
3.2.3 Contents	15

3.3	Giacomello’s DOOM Level Generator	18
3.3.1	Dataset Organization and Presentation	18
3.3.2	System Architecture	20
3.3.3	WAD Generation	21
3.3.4	Drawbacks	21
3.4	Improved Feature Map Extraction	23
4	Implementation	27
4.1	Dataset Procurement	27
4.2	System Architecture	28
4.3	Network Components	29
4.3.1	Topological Generator	29
4.3.2	Functional Generator	30
4.3.3	Data Flow	32
4.4	Model Optimization	33
4.4.1	Modified Loss Function	33
4.4.2	Training Algorithm	34
4.4.3	Updated WAD Generator	35
4.4.4	Experimental Setup	36
5	Results	39
5.1	Evaluation Metrics	39
5.2	Model Evaluation	41
5.2.1	Training Metrics	41
5.2.2	Proportions Analysis	46
5.2.3	Spatial Homogeneity	47
5.3	Sample Generation	48
5.3.1	Traditional WGAN-GP Model	48
5.3.2	Hybrid Model with Traditional cGAN	51
5.3.3	Hybrid Model with Modified cGAN	53
5.3.4	Condensed Category Map Comparison	56
6	Conclusion	57
6.1	Limitations	58
6.2	Future Works	58

A	Appendix A	63
B	Appendix B	67
	List of Figures	71
	List of Tables	73
	Acknowledgements	75

1 | Introduction

The perceived value of video games as a recreational activity is deeply rooted in the duration and intensity that they manage to keep players engaged. A few decades ago, as fewer games were released within any given period, developers would often try to prolong gameplay to validate the costs. This was achieved with irregular difficulty spikes and concealing information which would cause players to take several months to complete when coupled with the absence of the ability to retrieve one's progress. In comparison, despite numerous technological advances and man hours put into creating vast maps and longer stories, games today can be completed much sooner. This is partly due to adjustable difficulty, save points and widely available walkthroughs, thereby providing a shorter gameplay at a significant cost to the consumer. The novelty also quickly wears off on repeated playthroughs and feels monotonous, especially when there are limited to no variations in gameplay, affecting a game's ability to retain players engagement over longer duration.

Replayability is the player's desire to continue playing a game even after experiencing the content, either to optimize one's skill or to explore for its hidden content. It is beneficial as this kind of dedication arises from player satisfaction and has been shown to improve brand loyalty [7], yet many studios struggle to meet these expectations. This is usually attributed to the difficulty in maintaining player engagement over repeated sessions as it is critical that players don't know how every aspect of any level has been laid out. Some of the ways used to tackle this are by bringing in competitive aspects such as leaderboards or regular introduction of new content through updates which is heavily use of by many Massive Multiplayer Online Role Playing Games (MMORPGs) in order to maintain an active player base. The focus of this thesis is to address this problem by avoiding the monotony of traversing the same level through variations introduced in the level design.

1.1. Background

An approach that has long been resorted to combat repetitiveness in replay is procedural generation. Early adopters of the technology include games such as *Rogue* [36] with its procedurally generated dungeon and the trading game *Elite* [3], where each of its planets had randomly determined compositions. The former has spawned its own genre called ‘roguelike’ that initially used a similar tile-based procedural system. The notion of randomly generating levels solely utilizing existing assets at each instance appears more and more appealing as an alternative to constructing complex game worlds as it avoids excessive dependence on content design. This is especially true for newer games as there is greater demand for more content while maintaining graphics at a high fidelity. AAA titles such as *Civilization* [24] use random seeds in complex models to create different worlds for players to traverse and explore based on variation of terrain and resource placements while *No Man’s Sky* [8] progressively generates entire galaxies, including stars and planets with flora, fauna, and even alien encounters upon discovery.

1.2. Problem

While procedural generation is not new, it is an exhaustive process which requires consideration of minute details when building the system. AI has been developed for games, such as playing *Infinite Super Mario* levels [16] but have been less proficient when it comes to replicating creative tasks such as level design as the sample space natively adds a greater degree of complexity to the problem. In an attempt to avoid these pitfalls and consume vast amounts of resources for creating multiple levels, exploring deep learning seems to be a favorable course of action. There is a detailed method that is capable of generating DOOM level given certain input criteria through the use of conditional Wasserstein Generative Adversarial Networks with Gradient Penalty (WGAN-GP) [1] and the help of a large repository of custom-made levels created by the community to train the model. While the level layout generated by this model has proven adequate for a First Person Shooter (FPS), it is unable to place appropriate game objects into the level layout. This is because it is plagued with disproportionate generation of game object among the present categories and is biased with their placements as it clumps them together in certain sections of the level. This results in levels being unevenly saturated and predisposed towards certain types of game objects which hampers the level’s ability to properly captivate the player’s interest.

1.3. Goals

To improve the DOOM level generator by splitting the problems of planning the level layout and populating it with game objects through the introduction of a hybrid Generative Adversarial Network (GAN) architecture. This requires a repository of DOOM levels that need to be parsed into minimal structures that accurately represent the problematic features. With the feature maps serving as the examples to learn from, the hybrid architecture needs to be trained on the different aspects of level design by focusing on patterns present in specific feature maps that are sufficient to design a level until it can reproduce its effects with a degree of efficacy that is both topological and semantically sound. Once the networks accomplish their individual targets, it is compared with the prior design to extrapolate if there are any improvements with the added complexity before being implemented into a level generator. This entails that the generated set of features be reverse engineered into their respective game data to compose the encoded ‘lumps’ needed to serialize ‘.WAD’ files for the DOOM engine to correctly execute.

1.4. Thesis Structure

The thesis is composed of 6 chapters to compartmentalize the content relevant to this project into their respective sections. This is the 1st chapter which serves as a synopsis of the work that has been put into this thesis. The 2nd chapter showcases a brief overview of relevant research in related fields and suggests the theoretical knowledge required to construct this experiment. The 3rd chapter provides information pertaining to the dataset and its representation that will be used by the proposed model as well as the previous work that has been conducted regarding procedurally generating DOOM levels. The 4th chapter will deal with how the system is designed and implemented with the format with which the experiments are conducted. The 5th chapter illustrates the results that have been achieved with the designed system with regards to its predecessor through the defined metrics. The 6th chapter provides a summary of the work and preludes possible improvements that can be included in futures developments.

2 | Theory and Motivation

This chapter begins with a brief summary of works that are related to the topic of interest and a dive into the cutting edge of deep learning. This is followed by a review on past methods that have been researched and used for procedurally generating game levels. Finally, an introduction is provided regarding the theoretical foundations required to implement the techniques used in this thesis as well as an explanation outlining how they function.

2.1. Related Works

Artificial Intelligence in Video Games for the most part has referred to the behavior of non-playable characters, but it has been used for much more such as simulating the player or controlling in game systems [17, 22]. The Mario AI competitions [16] saw the mainstream adoption of Reinforcement Learning (RL) towards simulating the player's behavior with similar models outperforming human players in DOOM while also being able to navigate unknown maps [20]. On the other hand, Convolutional Neural Networks (CNN) have been used to predict the outcomes of FPS games to determine if the levels are biased towards any particular team through its maps or weapon parameters and accordingly modifies game data such as character stats and the level layout to control the flow of the game [17].

Through games, AIs have also been taught real world activities such as the city planner proposed to train agents to maximize the city population in the game SimCity using RL fractal networks which is further generalized to larger maps by means of a CNN with structured skip connections [5]. Mapping and texture design has also been widely researched with one such proposal entailing a two-stage GAN framework that can randomly generate height maps as well as infer texture maps from data provided by the NASA 'Visible Earth' project [2]. Its architecture consists of a Deep Convolutional GAN (DCGAN) used to generate the height maps with a Conditional Generative Adversarial Network (cGAN) for image to image translation to produce textures from the provided height maps.

2.2. State of the Art

Deep learning has come a long way since the inception of the perceptron in 1943 by McCulloch and Pitt [23]. There have been several architectural leaps that have provided it the maneuverability to permeate into a variety of problems that span multiple fields. This has never been more apparent than during the current resurgence following the pioneering of new techniques like transformers [39] that embed context into sparse representations of text to create 'Chat-GPT' and 'BERT', chatbots capable of conversations indistinguishable from humans on almost any topic. This comes after Recurrent Neural Networks such as the Long-Short Term Memory (LSTM) architecture [13] featured in language modeling and speech recognition until just a few years ago. These found it computationally infeasible to use large contexts from previous outputs, making subpar time-series predictions [44] that pale in comparison.

Larger strides have been made in the domain of images with GANs pioneering its own subset of image generative models such as the Cycle GAN [45], Projected GAN [31] and the Style GAN [18] among others, that have produced realistic visual data of buildings, terrain and textures [2, 38] as well as perform more precise tasks of modifying aspects of images to obtain specific results. Recent advances set forth by the advent of Diffusion Models have brought about systems such as 'DALL-E' and 'Stable Diffusion' that can capture patterns through sequentially denoising autoencoders [28], producing images using text-based cues at a quality never seen before. These have been used from reconstruction of medical images to the restoration and inpainting of video segments to reconstruct missing or damaged regions in an image. It has also been used to generate future frames by predicting the following image in a sequence through an inverse diffusion process [43].

These improvements have also been seen in the fields of computer vision with 2D and 3D detection capable of using a unified representation introduced by YOLOR [40], empowering systems by including implicit knowledge that serve various actions to allow them to learn multiple tasks such as object detection, multi-class labeling and feature embedding all at once. Graph based models have reached a level that can more directly impact medical applications with the predictions of drug interactions. The DSN-DDI [21] shows much promise as it has proven its ability to consistently predict interactions on even unseen drugs while also being able to combine drugs effectively by means of its dual view representation. It is capable of serving as a generalized framework for the discovery of novel drugs and alleviate resources wasted on unproductive clinical trials.

2.3. Literature Review on Procedural Generation of Game Levels

Hendrikx et al. provides a historical survey of methods that have been used for procedural generation for a variety of content in games, including isolated levels [12]. Some of the earlier games surveyed include NPPAngband, an Angband variant, which introduces fractal-based algorithms such as cellular automata for generating caves imitating indoor level design. Many of these algorithms consisted of generative grammars or pseudo random number generators combined with an advanced parameter space search such as genetic algorithms, often relying on a grid-based structures to generate these levels. Togelius et al. defines a taxonomy using evolutionary/stochastic based optimization algorithms and provides guidelines on how to approach the problem of representation and evaluation [35].

A personalized level generator was proposed by Summerville et al. using an LSTM architecture to classify training levels based on player paths extracted from video clips of them play [33]. The labeled structures serve as input to an auto-encoder which manufactures instances based on the learned design pattern. A similar architecture is used for rhythm-based games in Dance Dance Gradation to invent new levels with different degrees of difficulty based on acoustic features obtained from the audio track using a CNN [37]. Another approach that has been taken is training agents by modeling content generation as a Markov Decision Process to design levels of Zelda and Sokoban using various representations based on the autonomy the agent is given [19].

Park et al. generates levels with a tile-based system for Engage, an educational game that helps students learn computer science concepts and practices [25]. It uses a single encoded grid to represent the level layout and multi-step DCGANs to construct solvable levels. One performs training augmentation which creates more training data while another focuses on enhancing solvability by generating levels with higher solvability compared to those concocted by the former. Produced levels are judged using a shortest path algorithm for solvability and a tile comparison through k-nearest neighbors for novelty. It also introduces the concept of variable difficulty to the generator in its level design by classifying levels based on the degree of gaming skill needed to solve the level and its learning objectives to provide a personalized learning experience adapted to each individual player

2.4. Theoretical Background

2.4.1. Generative Adversarial Networks

GAN is a framework in which two neural networks are trained simultaneously in an adversarial setting. Generator G generates some target data during the training process while the Discriminator D helps it recover the data generating distribution. This works as a two-player min-max game [10] where the discriminator trains to classify whether the data x is from the real distribution p_{data} or not while the generator uses this feedback for the data it produces with noise z and trains to minimize the discriminator's ability to differentiate it from the real samples.

$$\min_{\{G\}} \max_{\{D\}} \mathbb{E}_{x \sim p_{data}(x)} [\log D(x)] + \mathbb{E}_{z \sim p_z(z)} [\log(1 - D(G(z)))] \quad (2.1)$$

The two networks are trained simultaneously through gradient descent performed over the loss which is calculated iteratively, then back-propagated over to the network's weights and biases until a local minimum is achieved. This is done by estimating lower-order moments \hat{m} and \hat{v} via an Adam optimizer that uses stochastic gradient descent with α as the step size and θ_0 being the initial parameter vector of the network.

$$\theta_t \leftarrow \theta_{t-1} - \alpha \cdot \hat{m}_t / (\sqrt{\hat{v}_t} + \epsilon) \quad (2.2)$$



Figure 2.1: Generated results using the MINST and TFD dataset from the work of Goodfellow et al.

2.4.2. Deep Convolution Generative Adversarial Networks

To overcome the difficulties of modifying GANs using the CNN architecture to work with images, DCGANs removes fully connected hidden layers for fractional-strided convolutions in the generator and strided convolutions in the discriminator [26]. This essentially swaps the optimization problem into learning an effective kernel that is able to generate the desired distribution. Spatial pooling functions are replaced since global average pooling increases model stability at the cost of convergence speed. To optimize this trade off, the convolutional features are directly connected to the input and output layer of both the networks. After each convolutional layer, Leaky Rectified Linear Unit (Leaky ReLU) activation functions are used in the generator and discriminator as it allows for back-propagation even for negative input values except for the output layer of the generator, which uses the Tanh function to quickly learn to saturate the color space for the training distribution completely.



Figure 2.2: Generated bedrooms using the LSUN dataset after five epochs of training from the work of Radford et al.

2.4.3. Wasserstien Generative Adversarial Networks with Gradient Penalty

As DCGANs use a Binary Cross Entropy Loss function instead of the KL divergence to minimize the Jensen Shannon divergence in GANs, they suffer from the generator's tendency to converge towards a singular solution to fool the discriminator (mode collapse) and vanishing gradients of discriminator predictions when trained to optimality. This is potentially from the generator's parameter space not being continuous and is addressed by using Wasserstein-1 distance to express the dissimilarity between two distributions. This provides the minimum cost to rearrange one distribution to another [1]. This Wasserstien GAN (WGAN) reconstructs the min-max game using the Kantorovich-Rubinstein duality with sample \tilde{x} belonging to the generated distribution \mathbb{P}_g and x to the real distribution \mathbb{P}_r .

$$\min_G \max_{D \in \mathcal{D}} \mathbb{E}_{x \sim \mathbb{P}_r} [D(x)] - \mathbb{E}_{\tilde{x} \sim \mathbb{P}_g} [D(\tilde{x})] \quad (2.3)$$

Another variation added to the traditional WGAN consists of a more robust loss function that accounts for higher order moments of the data distribution by introducing a penalty on the gradient norm for random samples \hat{x} . Unlike the k-Lipshitz constraint previously implemented through weight clipping, that biases the critic towards simpler functions, this gradient penalty enforces a 1-Lipshitz constraint on the gradient norm of the critic's output to its input [11].

$$L = \mathbb{E}_{\tilde{x} \sim \mathbb{P}_g} [D(\tilde{x})] - \mathbb{E}_{x \sim \mathbb{P}_r} [D(x)] + \lambda \mathbb{E}_{\tilde{x} \sim \mathbb{P}_{\tilde{x}}} [(\|\nabla_{\tilde{x}} D(\tilde{x})\|_2 - 1)^2] \quad (2.4)$$

DCGANs also normalize each feature independently across the batch to prevent the generator from collapsing all samples to a single point before every layer (except the generator's output layer and critic's input layer as it results in sample oscillation and model instability) but doing this changes the problem from individual mapping to mapping the entire batch for the discriminator. Thus, batch normalization is substituted with layer normalization in the critic since it does not introduce correlations between examples and the penalized training objective is now calculated using the norm of the critic's gradient with respect to each input independently.

2.4.4. U-Net

U-Nets are convolutional classification networks consisting of a contracting path to capture context followed by a symmetrically expanding path that enables localization [29] and concatenates the corresponding features from the contracting path. The network's loss is computed with a pixel-wise softmax through the use of a cross entropy loss function. The activation a_k for the set of classes K at the pixel position x is used to approximate a maximum function $p_k(x)$ that is meant to identify the true class for the image by bringing the results of the class with the maximum activation closer to 1 while dropping the output for the remainder of the classes near to 0.

$$p_k(x) = \exp(a_k(x)) / \left(\sum_{k'=1}^K \exp(a_{k'}(x)) \right) \quad (2.5)$$

$$E = \sum_{x \in \omega} w(x) \log(p_{\ell(x)}(x)) \quad (2.6)$$

The cross entropy then penalizes at each position the deviation of the activation $p_{\ell(x)}$ from the true label ℓ for each pixel and uses a pre-computed weight map w to give more importance to certain pixels, compensating for the different pixel frequencies in the dataset while the network is training. The initial parameter weights are drawn from a Gaussian distribution such that features have approximately unit variance. It is optimized through minimizing the Euclidean distance by averaging all plausible outputs with high momentum to ensure that the current optimization step is more based towards previously seen training samples.

2.4.5. Conditional Generative Adversarial Networks

Unlike conventional GANs, cGANs use a U-Net based architecture as its generator to perform image to image translations by mapping input images x towards the desired output images y using the noise z . It employs an $L1$ loss that can accurately capture low frequency information but produces blurry results since it cannot capture its high frequency sharpness. Thus, a Patch classifier is also incorporated which only penalizes structure at the scale of image patches to model high frequency structures.

$$L_{L1}(G) = \mathbb{E}_{x,y,z} [\|y - G(x, z)\|_1] \quad (2.7)$$

$$L_{cGAN}(G, D) = \mathbb{E}_y [\log D(x, y)] + \mathbb{E}_{x,z} [\log(1 - D(x, G(x, z)))] \quad (2.8)$$

This tackles the issues of blurriness by classifying patches of an image as real or generated and running the discriminator convolutionally across the image, averaging all its responses while assuming pixels separated by more than a set patch diameter are independent [15]. As such a network requires low-level information shared between the input and output, it uses small batches with high momentum while retaining batch normalization and dropout. The final objective is realized as a min-max game similar to that of the GAN with the inclusion of conditional inputs in both of its networks and an additional $L1$ loss that is amplified through λ .

$$\min_G \max_D L_{cGAN}(G, D) + \lambda L_{L1}(G) \quad (2.9)$$



Figure 2.3: Generated results on Satellite and translated images from Google Maps from the work of Isola et al.

3 | Representation and Prior Design

3.1. DOOM

In this thesis, the experiments are built around the game Doom [14] which was pivotal in the establishment of the FPS genre, putting forward novel techniques for 3D graphics as well as initiating support for custom modifications in its level format through its ports. It spawned dozens of copies, now called ‘DOOM Clones’ with Figure 3.1 showing that more than half the code base is cloned in certain older titles [4]. As one of the earliest 3D games, DOOM is a viable subject due to its many freely accessible custom levels that has been developed by the community over the years which are still tracked by websites such as ‘Doomworld’. It also benefits from its simpler 2D representation despite its 3D player environment, drastic simplifying the structures used to model and subsequently reducing the computational costs with its reduced sample space. This also benefits from widely available and detailed documentations, enabling it to be easily parsed and reverse engineered.




Figure 3.1: Percentage of Total Cloned Lines of Code with exact code clones (Type-1), similar code fragment with only different token names (Type-2) and similar algorithms (Type-3) from the work of Chen et al.

3.2. WAD

3.2.1. Format

Levels in DOOM are stored using the ‘.WAD’ format and contain all the necessary information for the DOOM engine to create the level [6]. The game was sold with 9 official levels which are Internal WADs or IWADs. These explicitly define each and every feature such as terrain and game objects as well as any graphics and sound required by the level. The dataset used in this experiment are solely Patch WADs or PWADs which are levels built by the community through map editors such as DOOM Builder and contain custom content or modifications that either add to or replace those already present in the originally referenced IWADs.

3.2.2. Structure

Each WAD is composed of byte data ordered as a sequence of lumps which is used to describe the feature to the DOOM Engine. It is preceded by a header that provides the WAD’s metadata and is succeeded by a directory containing metadata on every lump.

Section Length (bytes)	Section Name	Field Size (bytes)	Field Name	Description
12	Header	4	Identification	‘PWAD’ or ‘IWAD’ in ASCII
		4	Number of Lumps	Integer value of lumps included in WAD
		4	Table Offset	Integer pointer to resource dictionary
Variable	Lumps	-	Lump Data	Lumps stored as a byte stream
16 * Number of Lumps	Directory	4	Lump Position	Pointer to respective lump data
		4	Lump Size	Size of respective lump in bytes
		8	Lump Name	Lump Name in ASCII

Table 3.1: Sections of a WAD as seen in the work of Giacomello et al.

3.2.3. Contents

All the information that is used to define a level can be found within their respective lumps with each designated to a specific component of the level structure. The lumps defined in the Table 3.2 and Table 3.3 represent all of the mandatory lumps that should be present in every WAD to be deemed acceptable by the DOOM engine. Most lumps partition its entries into a set of fields with multiple entries in different lumps sometimes used to define a single entity. It also provides the necessary metadata that is used to signify special effects through fields allotted for flags that are present in all the unique lumps mentioned except for name and vertexes.

Lump Name	Description	Entry Size (bytes)	No. of Fields
NAME	The IWAD level name as an 'ExMy' label where 'x' is the episode number and 'y' is the mission number	2	1
THINGS	A list of every game object present that is not a wall, pavement or door, identified using an index with its position and orientation	10	5
LINEDEFS	A list of every line needed to signify all the walls, steps, and invisible boundaries such as event triggers that shape the level	14	7
SIDEDEFS	A list of referenced wall textures, specifying the sector that it faces and implicitly defining the sector's boundaries	30	6
VERTEXES	A list of x and y map coordinates that are referenced by linedef entries to create shapes on the map	4	2
SECTORS	A list of closed areas with the same floor and ceiling textures as well as height. Additional fields describe its lighting and the floor's special effects that usually damage the player	26	7

Table 3.2: Unique mandatory lumps in a WAD

The level building process begins through the designation of an individual sector and iteratively laying out its components until all the sectors are rendered. Sectors are constructed using the relevant linedefs which are obtained through associated sidedef entries that reference the respective sector. Linedefs are defined using two entries from the vertex lump and can either mention only a right sidedef for boundary walls or also a left sidedef when connecting two sector which can both be transparent to act as triggers using the corresponding flag .

There are certain constraints to the formation of lumps such as the presence of exactly one name lump in every WAD and the inability to render sectors above or below another since it uses a 2D map. Other functional lumps can be derived from the former using 3rd party software and serve to speed up the rendering process by avoiding runtime computation. All the lumps mentioned here except reject are mandatory for the engine to be able to build the level. These lumps include:

Lump Name	Description	Entry Size (bytes)	No. of Fields
SEGS	A list of wall segments comprising either a part or all of a linedef entry which are combined to form the various sub-sectors	12	6
SSECTORS	A list of sub-sectors which are the smallest undivided spaces incapable of obstructing the visibility of other walls within it	4	2
NODES	Branches in a 2D binary space partition tree sorting the view order of each sub-sector which speeds up the rendering process	28	8
BLOCKMAP	Pre-computed collision detection map without which objects and walls within the level cannot interact with any another	-	-
REJECT	Table to determining if an agent can ascertain the players location for visibility calculations used to optimize AI routines	1	8

Table 3.3: Derivable mandatory lumps in a WAD

These lumps speed up the rendering process by using a graph to represent the overall structure of the level and decide what needs to be shown by taking into account the player position and sorting the sub-sectors that are supposed to be visible to user. This is achieved by using the sub-sectors composed of segments as nodes in the graph that are rearranged every time the player moves through partition lines that divide the map and get iteratively smaller for each child until all sub-sectors have been encompassed.

Some of the derivable lumps have more unique structures such as the reject which expects byte sized entries with each bit representing visibility from a sector and the final byte padded to meet the size requirement. The blockmap on the other hand is preceded by a header of 8 bytes which indicates the coordinates of the origin and the number of blocks within the grid. This is followed by the offset of each linedef entry and the corresponding line segments.




Figure 3.2: Sample DOOM PWAD visualized using DOOM Builder where orange indicates enemy positions, yellow represents pick-ups such as power-ups/ammunition and green is the starting position of the players

3.3. Giacomello’s DOOM Level Generator

3.3.1. Dataset Organization and Presentation

WADs are converted into feature maps which function as the training data for the neural network. This is done by translating extracted byte data by means of a custom parser and composing them into local data structures for the aforementioned lumps. By focusing on the minimum requirements necessary for creating a DOOM level, the parsed dataset is limited to 4 feature maps which represent the level topology as well as its functional components. To this end, the wall, floor, and its height proved capable of accurately representing the level layout while the things map handles the presence of gameplay elements. After parsing the WAD, relevant coordinates for each map are marked on separate matrices. The dimensions of this matrix are allocated based on the length and width of the given map, scaled down from its original size by the radius of the smallest function object, in this case by 32 times.

The level shape is sketched out via sectors through finding all sidedefs with the same number in the sector field and obtaining all the linedefs that reference those sidedefs. The retrieved list of linedefs is used to identify the pertinent vertices which are then adjusted for the new scale of the map. The wall map is outlined by drawing lines between the referred vertices for each linedef only if it contains at least one sidedef that is not transparent. The height map is computed sector wise by extrapolating the perimeter given by these lines and iteratively traces out the respective polygons that envelope the area containing that sector. The floor map is obtained by creating a mask of the positive value in the height map while the things map uses the scaled coordinates to encode pixels with values from the things type dictionary.

A further set of scalar features such as the ratio between the level area and its convex hull, mean room size, largest room size, room count, diameter of circle with the same level area, length of the longest and shortest axis of the level as well as the skewness and peak sharpness in the pixel distance distribution from its closest wall are also computed and appended to the maps as part of the dataset. This is done in order to provide some degree of control over the generation process as they are easily extracted and impervious to noise. The dataset is then stored as Tensorflow records to avoid recalculation.




Figure 3.3: Floor map, Height map, Wall map, and Things map (from left to right) of a sample level

The feature maps shown in Figure 3.3 represent the level in its bare minimum and are used to derive all the mandatory lumps required to furnish a level while delivering the gameplay envisaged by the genre. The features maps are engineered as such:

- FLOOR MAP: Depicts the floor layout of the level by segregating the available space into the portion that lies within bounds of the level using pixel values of 255 to denote them and leaves that which remain with a pixel value of 0.
- HEIGHT MAP: Marks the height of various sections of the map by encoding pixel values in gradients for segregating each sector with a maximum value of 255 and those outside the bounds of the floor map as 0.
- WALL MAP: Compartmentalizes the level into separate rooms and visualizes the points of access for each of the rooms by tracing the positions of the walls used to partition it. Walls are marked with a pixel value of 255 while the rest are left with a pixel value of 0.
- THINGS MAP: Indicates the position of the game objects on the map by identifying the type of object through pixel values assigned using the id allotted to it in the things type dictionary while the absence of objects is represented as 0.

3.3.2. System Architecture




Figure 3.4: Conditional WGAN-GP architecture used for the generation DOOM levels from the work of Giacomello et al.

a sigmoid activation function as it better coincides with the target format. It uses the feature vectors accompanied by a gaussian noise vector as input and learns its distribution while the critic uses the real feature maps and the generated ones to grasp the differences between them given those feature vectors. The model's predictions are then evaluated based on the Wasserstien distance from its expected value. The parameters are then adjusted using an Adam Optimizer with the generated levels compared using metrics such as corner count, structural similarity index and encoding error to assess the quality of the developed architecture.

Giacomello approaches the problem of procedurally generating DOOM levels by means of a Conditional Wassertien GAN with a Gradient Penalty [9]. By rescaling the feature maps extracted from the TFrecords into the same shape of 128 by 128 pixels, the generator is trained on a normalized version of the dataset which is further enlarged by including three more orientations by rotating the maps in multiples of 90-degree. The extracted input scalars called 'Feature Vectors' are also appended to the feature maps to correlate the parameter to their respective level.

The generative network has five hidden layers which comprise of one linear layer followed by convolutional layers with a final convolutional output layer applying

3.3.3. WAD Generation

Once the network is capable of generating maps recognizable as DOOM levels, a custom WAD writer is used to reverse engineer the matrices into lumps of byte data representing the mandatory components of a PWAD. By using the floor, wall, height, and things feature maps, it prepares the level layout by reconstructing the information into an undirected weighted graph. This is done through computing the respective room maps by means of a segmentation algorithm to compartmentalize sections of the walkable area of the level. The graph is then constructed outwards from the center and is decorated with various features while keeping the textures constant.

Level paths that can be traversed are collected using a minimum spanning trees algorithm on the graph from the section of the level with the least entrances which is used as the starting position of the level. Based on the computed solutions, an exit trigger is generated at the end of the longest floor path. Once the graph is complete, its nodes are translated into their respective lumps by scaling up each of the coordinates by 64 before it is finally committed with a level name and generated as a WAD. Four texture entries are required in all such PWADs for DOOM to run properly which are '*FLOOR4_8*', '*SFLR6_1*', '*MFLR8_4*', and '*FLOOR7_2*'. The first three are needed as backgrounds for the episode end texts. The last is what is shown outside the border of the display window when not being run in full screen.

3.3.4. Drawbacks

Due to the use of interpolation techniques to resize the image into identical sizes, the information within the images is not exactly translated when performing such transformations. Effects of distortion and blurring can be addressed or ignored during the dataset preparation for the topological maps as the variation is consistent in the height map while the floor and wall maps are indifferent to exact pixel values. Sparse images such as the things map that depend on the pixel encoded values on the other hand are contorted beyond permissible levels of tolerance, making it to be difficult to maintain its integrity. In the case of the things map, this has the effect of pixels taking on the value of other objects, coercing them towards objects with lower assigned ids. Multiple objects close together also get overlooked when larger maps are scaled down from their original size based on the order the object is presented with the WAD. This is due to compacting the level space which results in large units of area converging into a single pixel and the indiscriminate superseding by the last object in the list regardless of category, further dissociating the relations between object types.




Figure 3.5: Height and Things maps (From left to right) of the sample levels (above) compared to its resized counterpart (below)

Maps smaller than the target size of 128 by 128 pixels meanwhile must be enlarged to compensate for the lack of area which leads to objects being represented through the multiple surrounding pixels, further deteriorating the quality of the map. It also leads to large empty spaces throughout the level which can throw off the network's ability to learn and populate the level space completely.

There are bugs encountered in the generation process as disconnected sections of the level are inaccessible and the objects present within these portions of the level do not meaningfully contribute in any way to the actual proportions of the category types. If the player spawns in a section that has a greater population of monsters while another detached section has more ammunition, it presents the player with little to no opportunity to exploit the totality of resources available. Attached segments can also be inaccessible when the connecting corridors are too small to traverse, only allowing the player to peek through which is detrimental to gameplay.

3.4. Improved Feature Map Extraction

To alleviate some of the problems encountered, modifications have been made to Giacomello's WAD Editor to parse the WADs for a better representation of the sufficient subset of unique feature used to design a level. From Figure 3.6, it can be noted that most maps lie below 15000 units in both dimensions from the figure. The parsed maps are now scaled down to 256 by 256 pixels with levels below 15000 units being relatively padded by two row or columns for every 300 units smaller than the mentioned limit in its respective dimension to preserve objects density. Although there is still some overlap among game objects, integrity of pixel values is ensured by using only coordinate rescaling instead of resizing throughout the entire process.




Figure 3.6: Area distribution of the filtered dataset

While retaining the topological maps, the things map has been further divided into their individual category maps to provide a more coherent context for the system to process. Objects of the respective category are designated an id, acquired by arranging them in

the order of significance of the effect it has on gameplay. Only functional elements that are required to fully experience the game have been factored in and can be recombined to recreate the thingsmap. The categories included are monsters, ammunition, power-ups, artifacts, and weapons.




Figure 3.7: Floor, Height, Wall, Monsters, Ammunition, Powerups, Artifacts and Weapons map (From left to right, top to bottom) of the sample level using coordinate rescaling

The newly added set of feature maps apart from the included floor, height and wall maps from the previous design are:

- MONSTERS MAP: Assigns the respective positions of enemies with the pixel value corresponding to its monster id and 0 representing its absence.
- AMMUNITIONS MAP: Assigns the respective positions of ammunition used by the various weapons with the pixel value corresponding to its ammunition id and 0 representing its absence.
- POWERUPS MAP: Assigns the respective positions of armor and first aid with the pixel value corresponding to its powerup id and 0 representing its absence.
- ARTIFACTS MAP: Assigns the respective positions of specialized abilities with the pixel value corresponding to its artifact id and 0 representing its absence.
- WEAPONS MAP: Assigns the respective positions of the various weapons with the pixel value corresponding to its weapon id and 0 representing its absence.

Mentioned in Table 3.4 and Table 3.5 are the list of game objects that have been taken into consideration as they provide a functional value to the setting of the game. Apart from the start position of the player, the remainder of the objects are generated using the implemented model.

Id	Things Type Id	Description	Category	Category Id
1	0	Empty	-	0
2	3	Player 1 Start	other	1
3	75	Arachnotron	monster	11
4	76	Arch-Vile	monster	14
5	77	Baron of Hell	monster	16
6	78	Cacodemon	monster	9
7	79	Chaingunner	monster	7
8	81	Cyberdemon	monster	17
9	82	Demon	monster	4
10	83	Former Human Trooper	monster	3
11	84	Former Human Sergeant	monster	5
12	85	Hell Knight	monster	15
13	86	Imp	monster	1
14	87	Lost Soul	monster	8
15	88	Mancubus	monster	13
16	89	Pain Elemental	monster	10
17	90	Revenant	monster	12
18	91	Spectre	monster	6
19	92	Spider Mastermind	monster	18
20	93	Wolfenstein SS	monster	2
21	94	Ammo Clip	ammunitions	2
22	95	Box of ammo	ammunitions	6
23	96	Box of rockets	ammunitions	7
24	97	Box of shells	ammunitions	5
25	98	Cell charge	ammunitions	4

Table 3.4: List of relevant game objects with category ids (1 of 2)

Id	Things Type Id	Description	Category	Category Id
25	98	Cell charge	ammunitions	4
26	99	Cell charge pack	ammunitions	8
27	100	Rocket	ammunitions	3
28	101	Shotgun shells	ammunitions	1
29	102	BFG9000	weapons	7
30	103	Chaingun	weapons	3
31	104	Chainsaw	weapons	1
32	105	Plasma Rifle	weapons	6
33	106	Rocket Launcher	weapons	5
34	107	Shotgun	weapons	2
35	108	Super shotgun	weapons	4
36	109	Backpack	powerups	3
37	110	Blue armor	powerups	5
38	111	Green armor	powerups	2
39	112	Medikit	powerups	4
40	113	Radiation suit	powerups	6
41	114	Stimpack	powerups	1
42	115	Berserk	artifacts	3
43	116	Computer map	artifacts	5
44	117	Health potion	artifacts	1
45	118	Invisibility	artifacts	8
46	119	Invulnerability	artifacts	9
47	120	Light amplification visor	artifacts	7
48	121	Megasphere	artifacts	6
49	122	Soul sphere	artifacts	4
50	123	Spiritual armor	artifacts	2

Table 3.5: List of relevant game objects with category ids (2 of 2)

4 | Implementation

4.1. Dataset Procurement

The data used to train and validate the models is obtained by scraping custom made doom levels from 'doomworld.com', a source mirror of Idgames Archive. A scraper is made to access the domain and iterate over the web pages of the available levels within each subcategory present excluding 'Deathmatch', 'Megawads' and 'Ports'. This choice has been made to avoid mixing different types of levels as player vs player environments are structurally different from player vs enemy environments. Ports has been ignored because it requires modifying the game engine to run the level which is complicated to characterize in a format decipherable by the model and stably implement into PWADs. The web scraper then downloads WADs and stores relevant information in a JSON file which doubles as a register to catalog the collection at hand.

Among the custom levels present in Doomworld's repository, 1969 of the available levels do not fall within the excluded categories mentioned above. In an effort to reduce the complexity of the problem, WADs containing only a single level have been included to ease the requirements of the system with a simpler manner of representation. Similarly, constraints have been placed on the levels to ensure its faithfulness to the genre is respected such as accepting only those that contain at least one monster and weapon to improve the overall consistency within the dataset. Finally, to forgo the intricacies that arises with teleporters, levels that have large, detached segments have not been included, leaving only the remaining 1279 levels which are taken into consideration and included in the finalized dataset. The WADs that have been determined to meet the desired criteria are then parsed and have its feature maps extracted to serve as the training data for the models.

4.2. System Architecture

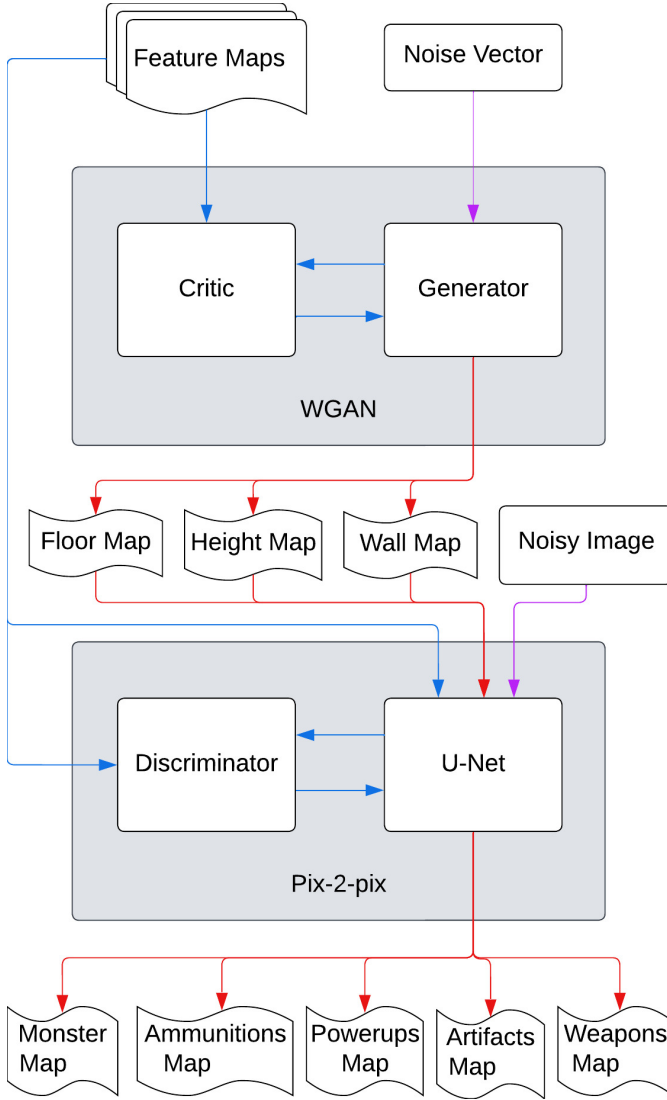


Figure 4.1: The hybrid GAN architecture where blue lines represent the training process, red line are used in the generation process and the purple lines are for both

The proposed system is a hybrid architecture consisting of a WGAN-GP and a cGAN model that are trained on specific feature maps to grasp particular aspects of level design. The WGAN-GP is provided with those pertaining to the level's topology to learn to draft coherent level layouts while the cGAN is taught the functional aspects of the level and uses it to spawn game objects in a manner that positively influences the game dynamics within the generated level. These networks are built using the Keras Sequential API from the Tensorflow libraries in Python 3.9.

The models are split into individual networks and are trained separately due to the size restriction on the memory of the GPU. During the generation process, the networks are run consecutively with the WGAN-GP executed first so that its results can be fed into the cGAN after certain post processing. The results obtained from both the networks are then concatenated together to attain a complete set of feature maps which are then used to manufacture the mandatory lumps and assembled into

a WAD. The remaining lumps are procured using a 3rd party software and this process is explained in greater detail in the following sections.

4.3. Network Components

4.3.1. Topological Generator

The subsystem is designed using Radford and Metz's Deep Convolutional Generative Adversarial Network architecture with Wasserstien distance used to calculate the loss as undertaken by Arjovsky et al. The generator is trained alongside a critic implementing alterations by Gulrajani et al. which incorporates a gradient penalty to its loss estimation and has the inclusion of additional training steps compared to that of the generator. The aim is to teach this module the topological structure of levels which is defined by its floor space, the enclosing wall layout, and the respective height of each section.

The data handled involves a random gaussian noise vector Z , the topological feature maps X and the critic's classification of the allocated map Y . The topological feature maps consist of the floormap, wallmap and heightmap images and those obtained from the scraped dataset represents the real distribution X_{True} that the generator is to learn. These are used in combination with images from the generated distribution to help the critic determines which features needs to be extracted to correctly classify a provided image. The sampled feature maps X_{Gen} from the generated distribution is procured using the generator G through manipulation of the noise vector i.e.,

$$X_{Gen} = G(Z) \quad (4.1)$$




Figure 4.2: Topological generator of the hybrid architecture

The critic C uses an identical architecture in the same manner with the only difference being that it operates on three feature maps. It processes the maps with its hidden layers progressively through repetitive steps of down convolution followed by layer normalization and Leaky ReLU activation until it is finally condensed into a single output assigning the map belonging either to the real distribution Y_{True} and the generated distribution Y_{Gen} to a value between 0 and 1. It is equivalent to the logits functions where it optimizes the critic's prediction as follows:

- $Y_{Gen} = C(X_{Gen}) = \text{Logits}(X_{Gen})$ are brought closer to 0 as it indicates that the sample is generated.
- $Y_{True} = C(X_{True}) = \text{Logits}(X_{True})$ are brought closer to 1 which symbolizes that the sample is real.

Similarly the generator is implemented without the scalar features, taking in a 100-element noise vector as input and manipulating the tensor length through a dense hidden layer to project the noise into a larger tensor. The tensor is then reshaped into a set of 8 by 8 images which are then expanded through a series of up convolutions until the final output layer produces a floor, wall, and height maps for the generated level of pixel size 256 by 256. A sigmoid activation function is used in the output layer instead of the orthodox tanh function as it is more appropriate since the constraints the feature maps lie between have been normalized.

4.3.2. Functional Generator

The module is based on Isola et al.'s Conditional Adversarial Network with modifications made to its loss function to help better extract the target distribution in the given circumstance. The generator's architecture uses the three 256 by 256 pixel topological maps X_I with an image of gaussian noise Z of the same size and progressively transforms the tensors into the category maps X_T consisting of the monsters, ammunition, powerups, artifacts and weapons maps. The network is initially trained to generate category maps X_{T_Gen} using the sample input X_{I_True} from the real distribution. It is then compared with the original object map X_{T_True} by the discriminator. During the generation phase, the network uses the maps X_{I_Gen} produced by the topological generator to obtain the newly generated category maps X_{T_GGen} .

$$X_{T_Gen} = G(Z|X_{I_True}) \quad (4.2)$$

$$X_{T_GGen} = G(Z|X_{I_Gen}) \quad (4.3)$$

The generator is trained in a method which requires it to retain its batch normalization even after the training process. Instead of using a tanh activation function for the generator output layer, a ReLU activation function has been used as it has proved to provide better image translation. The generator loss function has also been modified to account for the problem's parameters and as such it contains the adversarial loss i.e binary cross entropy but has the $L1$ loss substituted with the and mask loss which is explained in greater detail in the next section.

The discriminator classifies the target maps with its input map samples from both the real and generated sample distributions much like that of the WGAN-GP with the only differences being that it is provided with all mentioned feature maps and that it uses a Binary Cross Entropy loss function. Both the generator and discriminator use modules of the form of a convolution, followed by batch normalization and a ReLU activation when upscaling or a Leaky ReLU while downscaling. The predictions are composed as patches which are calculated similar to logits and are optimized as

- $Y_{Gen} = D(X_{I_True} + X_{T_Gen}) = \text{Logits}(X_{I_True} + X_{T_Gen})$ are brought closer to 0 as it indicates that the sample is generated.
- $Y_{True} = D(X_{I_True} + X_{T_True}) = \text{Logits}(X_{I_True} + X_{T_True})$ are brought closer to 1 which symbolizes that the sample is real.

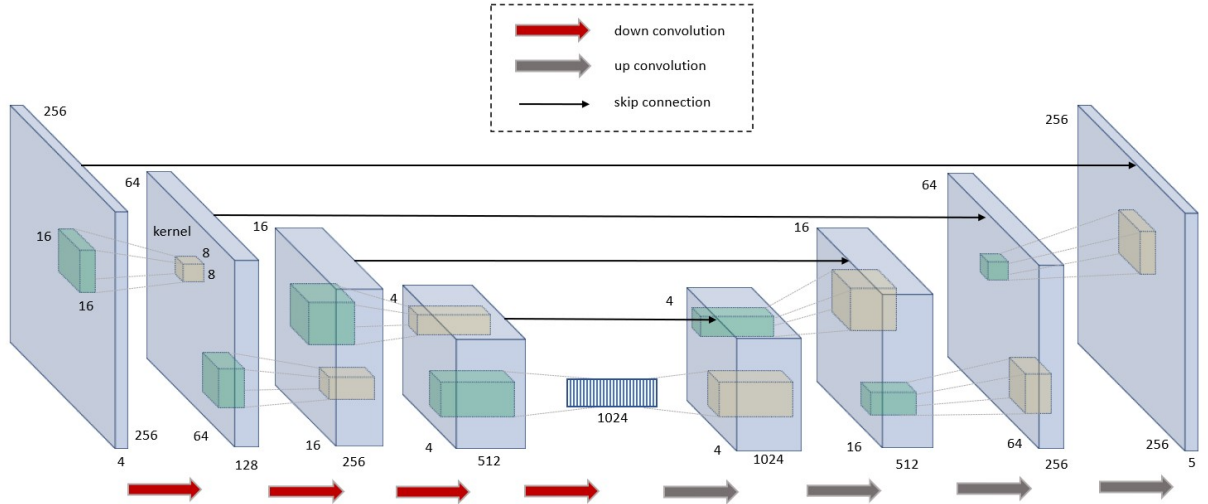


Figure 4.3: Functional generator of the hybrid architecture

4.3.3. Data Flow




Figure 4.4: Data flow diagram

The flow of data varies between the two use cases that are being performed by the system. The training case is broken into 2 steps, the first of which is the dataset processing phase which begins with the dataset procurement of level repository from the archives that appears in the top left of Figure 4.4 to obtain the dataset as TFRecords. Once the dataset is procured, the model optimization phase uses the dataset until the constructed models have been sufficiently taught to obtain the feature maps in a manner that can mimic the dataset.

The next case consists of the generation phase, during which the models transform random noise into feature maps with the WGAN-GP's outputs go through a post processing phase to help clear noise and artifacts. The resultant maps are then validated to ensure the absence of anomalies before being supplied to the cGAN. Then the model outputs are tested through evaluation metrics to inform of its quality and converted into WADs for examination of possible in gameplay.

4.4. Model Optimization

4.4.1. Modified Loss Function

While the WGAN-GP uses the loss function defined in Section 2.4.3, the cGAN network performs inconsistently with the $L1$ loss on different optimization attempts from the random shuffling of the training dataset. Thus, a suitable loss function has been hand crafted for the problem which moves away from the $L1$ loss by replacing it with more milder constraints. The first of which is the mask loss L_M which is used to compute the number of objects that deviate from the bounds of the level.

$$L_M = \ln|1 + \sum_{i=1}^n \sum_{j=1}^n b_{ij} \cdot B(X_{T_Gen_{ij}})| \quad (4.4)$$

Where b indicates if pixel ij is not within the level bounds and B is the binary of the presence of a game object at that pixel of the generated maps. The other is the object loss L_o which is used to approximate the deviation of the generated objects from the true proportions of the categories in that of the true target maps of the level. It is obtained through finding the absolute sum error ϵ_{as} of the generated feature maps from that of the target feature maps for each category k .

$$\epsilon_{as^k} = |\sum_{i=1}^n \sum_{j=1}^n X_{T_True_{ij}^k} - X_{T_Gen_{ij}^k}| \quad (4.5)$$

$$L_O = \frac{1}{|K|} \sum_{k=1}^K \lambda \cdot \ln|1 + \frac{\epsilon_{as^k}}{1 + \sum_{i=1}^n \sum_{j=1}^n X_{T_True_{ij}^k}}| \quad (4.6)$$

This error is taken in proportion to the number of target objects to ensure that the categories with fewer objects do not get ignored. Just as the mask loss, a natural log is used to modulate large loss values towards a more moderate scaling with λ amplify its significance. Finally, the mean of this loss is taken across the available categories K is then summed with the GAN loss from Equation 2.8 and the mask loss to derive the total generator's loss.

$$L_G = -\mathbb{E}_{x,z}[\log(D(X_{I_True}, X_{T_Gen}))] + L_M + L_O \quad (4.7)$$

4.4.2. Training Algorithm

The system is trained through iteratively executing the training operations over 1023 training samples X_{train} with the models validated over the remaining 256 samples X_{Val} . At the end of each epoch, reference samples are computed using the same seed of gaussian noise with a mean of 0 and standard deviation of 1. These are used to provide a visual understanding as to the development of the model over the training cycle. As the architecture is composed of two separate models, the training algorithm has been distinctly tabulated with one detailing the processes performed by the WGAN-GP networks and another for the cGAN networks.

This dataset has been artificially enlarged for the WGAN-GP networks to 8064 and 2048 for training and validation respectively through the consideration of the three 90-degree rotations and its mirror images to exploit for the differences in level orientation. Each of the WGAN-GP networks has been trained for 100 epochs using a batch size of 16 with validation performed every 100 batches. The critic is trained for 3 steps for every step that the generator is trained instead of the suggested 5 but performs similarly when compared. A weight of 10 is multiplied to the gradient penalty for the critic loss. Both networks utilize Adam optimizers with a learning rate of 2e-4, beta decay of 0.5 and momentum set at 0.999.

Process	Inputs	Outputs	Periodicity	Evaluated Opertors
Train G	$Z_U[0, 1]$	L_{G_Train}	3	X_{Gen}, Y_{Gen}
Train C	$X_{Train}, Z_U[0, 1]$	L_{C_Train}	1	$X_{Gen}, Y_{Gen}, Y_{Train}$
Validation	$X_{Val}, Z_U[0, 1]$	L_{G_Val}, L_{C_Val}	300	$X_{Gen}, Y_{Gen}, Y_{Val}$
Reference Sample	Z_{Ref}	X_{G_Ref}	1512	-
Network Checkpoint	G, G_{Optim}	checkpoint	1512	-

Table 4.1: Training operation of WGAN-GP models

The hybrid network on the other hand does not have its dataset enlarged since they are trained with a batch size of 1 as consecutive samples with altered orientations causes it to retain unnecessary bias when presented with unusual samples from the dataset. It is also trained similarly to the WGAN-GP, using the same hyperparameters with the exception being that the Adam optimizer has a learning rate of 6e-5 to reduce the impact of each batch and avoid reaching a local minimum and a weight of 10 for the object loss.

Process	Inputs	Outputs	Periodicity	Evaluated Opertors
Train G	$X_{I_Train}, Z_U[0, 1]$	L_{G_Train}	3	X_{T_Gen}, Y_{Gen}
Train C	$X_{Train}, Z_U[0, 1]$	L_{D_Train}	1	$X_{T_Gen}, Y_{Gen}, Y_{Train}$
Validation	$X_{I_Val}, X_{T_Val}, Z_U[0, 1]$	L_{G_Val}, L_{D_Val}	100	$X_{T_Gen}, Y_{Gen}, Y_{Val}$
Metrics	X_{I_Val}, X_{T_Gen}	$metric_{eval}$	100	-
Reference Sample	X_{I_ref}, Z_{Ref}	X_{G_Ref}	1023	-
Network Checkpoint	G, G_{Optim}	checkpoint	1023	-

Table 4.2: Training operation of cGAN models

4.4.3. Updated WAD Generator

Feature maps are reverse engineered into the mandatory lumps necessary for the DOOM engine with the help of certain refinements made to Giacomello’s WAD Editor. Firstly, detached regions of the level have been addressed through the introduction of post processing before generation of category maps. It removes holes and small objects from the floor map to ensure that it is fully connected before generating the corresponding category maps. Another issue is the presence of inaccessible paths that are caused by narrow corridors between the sections. This happens when the scale size of pixel wide sectors is smaller than the player. It has been tackled by ensuring that the minimum width of paths represented by pixel is greater than the dimensions of the character by enforcing a scaling factor of 128 while never exceeding the bounds of the available level space.

The Doom Engine also uses a Binary Space Partitioning Algorithm for pre-computing the Hidden Surface Determination or occlusion culling. In this work, the tool ZenNode [30] has been used in the last stage of the pipeline, to produce playable DOOM levels from the network output. ZenNode is a 3rd party software used to generate lumps that does not directly affect the level design. It has additional functionalities with a customizable Binary Space Partition to ignore certain lindefs and specify unique sectors, though its main function is to provide incomplete PWADs with blockmap and reject resources automatically, given that the remainder of the lumps have already been generated using the feature maps.

4.4.4. Experimental Setup

The experiment performed is a comparative analysis between the proposed and prior architectures with the traditional model added in to provide plausible justification for the added modifications to the hybrid network. While training the networks, similar models have been constructed near indistinguishable to minimize any variations that are not due to the inherent differences between the models. To this end, the networks use the same number of layers, hyperparameters, kernel size and stride wherever possible.

The experiment is set up initially to compare the cGAN networks through metrics calculated during the training phase. Although it may seem tempting to simplify the functional generator by means of using just a U-Net architecture which only implements the object and mask loss as discussed in the previous section, it is incapable of mimicking the subtle variations of object placements through performing the image-to-image translation purely with a fixed loss function and hence has been rejected for this system.

The traditional WGAN-GP network is not taken into account for the training metrics as it does not use a two-step process making it difficult to calculate them without context. It also does not use the aforementioned category maps due to its generator's tendency towards producing predominantly zeros matrices for those features due to the sparse nature of the distribution. This architecture is found to at least be able to generate some game objects when trained with the thingsmap which has been included in the post training evaluation of the networks. Post training evaluation is done over 100 samples from the real and generated datasets.

Model	G_{Loss} Function	Feature Maps	G Convolutional Layers (filters)	D Convolutional Layers (filters)
WGAN-GP	Wasserstien Distance with Gradient Penalty	floormap, heightmap, wallmap, thingsmap	4 (1024, 512, 256, 128)	4 (128, 256, 512, 1024)
WGAN-GP + Tradition cGAN	Wasserstien Distance with Gradient Penalty	floormap, heightmap, wallmap	4 (1024, 512, 256, 128)	4 (128, 256, 512, 1024)
	Cross Entropy + $L1$ Norm	monstersmap, ammunitionsmap, powerupsmap, artifactsmap, weaponsmap	7 (128, 256, 512, 1024, 1024, 512, 256)	4 (128, 256, 512, 1024)
WGAN-GP + Modified cGAN	Wasserstien Distance with Gradient Penalty	floormap, heightmap, wallmap	4 (1024, 512, 256, 128)	4 (128, 256, 512, 1024)
	Cross Entropy + Object Loss + Mask Loss	monstersmap, ammunitionsmap, powerupsmap, artifactsmap, weaponsmap	7 (128, 256, 512, 1024, 1024, 512, 256)	4 (128, 256, 512, 1024)

Table 4.3: Trained architectures used for the comparative study

5 | Results

5.1. Evaluation Metrics

The focus of this portion of the thesis is to quantify the evaluations that have been made through the results obtained from the experiments using the proposed metrics. The preliminary analysis is done to provide plausible justification for the use of the modified cGAN network through the training metrics that have been acquired alongside its validation by showcasing their differences. The metrics consist of the discriminator's ability to classify the levels, the unique information needed to describe the level as well as the following statistics:

- **Mean Object Count** – It is the count of all objects within a level by the total area that the level occupies. It provides an approximate object density that is found in levels and serves as a constraint to check if the object maps are leaning towards over saturation or under saturation as the model must learn to balance increasing the availability of objects while avoiding overpopulating a level by making it too congested.
- **Mean Encoding Error** – The network generates images where pixels are encoded using continuous values which can be rounded towards the closest object type. As the networks use a ReLU activation function in the output layer due to the tendency of the cGANs towards predominately empty images with sigmoid activation, there is the ability to encode values beyond the limits of assigned ids of the respective categories. This error is used to measure the inaccuracy in pixel values which would overflow into other categories if not addressed.
- **Mean Out of bounds Error** – When objects are generated outside the bounds of the level, they are ignored in the level generation process as they are inherently inaccessible to the player and thus not part of the actual level. This error is used to quantify how often this occurs and the degree to which the network is unable to comprehend the constraints within which game objects need to be placed.

It is difficult to compute the former metrics during the training of the traditional WGAN-GP since it is learning the complete distribution and they require a definitive level structure which is not yet established, making it impractical to compare their respective graphs. Instead a set of post training metrics are used to provide a more appropriate analysis. These include the performance of the generator in understanding the problem by identifying visual artifacts in the presentation of the things map and distinguishing the samples intuitively. as well as the analyzing the object distribution, both spatially and categorically. While not it is not completely comprehensive, it should provide conclusive evidence of their abilities. These later techniques assess the presence of a game object around the surrounding context, be it the local section or the encompassing level and are described as follows:

- Proportions Analysis – It is based on the proportions of the various categories based on the types of objects that it uses to populate the level. This is used to understand how well the model is able to grasp the relation between the requirement of placing objects from different categories and significance of their individual types. This in junction with the count of the individual object type from the available categories provide a clear picture on the biases and tendencies of the generator when applying a positive encoding on any pixel.
- RipleyK Function – It is used to measure the spatial distribution of the generated objects [27]. It determines whether points have a random, dispersed, or clustered distribution pattern. The function is calculated by taking the average occurrence of events within a predetermined radius of every point iteratively which provides an advantage over other methods since it is independent of the density of the point pattern and thus can describe the spatial point patterns at many different scales [42]. The function is evaluated as follows:

$$K(s) = \lambda^{-1}n^{-1} \sum_{i=1}^n \sum_{j \neq 1}^n I(d_{ij} < s) \quad (5.1)$$

where d_{ij} represents the distance between points i and j , s is the threshold radius within which the point is considered and n is the number of points that are sampled from the distribution with λ being the reduction factor of the function.

5.2. Model Evaluation

5.2.1. Training Metrics

By using the loss computed during the training sequence, we are able to compare the variation in the discriminator's ability to identify the generated levels. As both cGANs use different loss functions in their generators, it is difficult to directly compare their assessments and thus more emphasis will be given between graphs of the discriminators from both models of the hybrid architecture. Although validation data is not as useful in GANs to assess overfitting and/or implement early stopping, it is still done to hold back some real samples for testing the discriminating network to compare how well it can classify images among the functional generators [34]. This can be seen in Figure 5.1 with the discriminator validation loss of the traditional cGAN rising above that of its training counterpart around the 40th epoch while the network is only able to generate viable samples after the 80th epoch.




Figure 5.1: Discriminator Loss of the traditional cGAN model




Figure 5.2: Discriminator Loss of the modified cGAN model

As seen in Figure 5.1 and Figure 5.2, the traditional cGAN is better at deceiving its discriminator with a smaller loss value but is plagued with inconsistencies due to the stricter constraints of the $L1$ norm present in the generator loss when handling outliers in the training dataset, leading to a more noisier graph when compared to the laxer restrictions imposed by the object and mask loss. This is also seen in the fact that the traditional cGAN performs very differently from independent instances of the complete training cycle and seems to be extremely sensitive towards the shuffling process while the modified cGAN is able to repeatedly settle at a minima that provides consistent results.




Figure 5.3: Mean entropy of the generated category maps




Figure 5.4: Mean object count per unit area of the generated category maps

Figure 5.3 and Figure 5.4 show the similarities between the cGAN networks with amount of randomness and the object counts being approximately the same during the training process which is misleading when not taken into context with the errors produced by the networks. The next set of metrics provide a better understanding of this inconsistency of the traditional cGAN with the encoding error showing lesser accuracy for assigning values within the limits of the object ids and the out of bounds error highlighting the positioning the objects outside the prescribed level space




Figure 5.5: Mean encoding error of the generated category maps

Even though the traditional cGAN provides a object count in line with its modified counterpart, most of the objects that have been generated are not within the limit of the dictionary ids as seen in Figure 5.11 and Figure 5.12 which has an negative effect to the otherwise normal object count after the correction is made. This is caused by outliers that occur within the repository of feature maps because some custom levels having an unusual distribution such as higher count of more significant object types in some categories while the others only having a few less significant object types.




Figure 5.6: Mean out of bounds error of the generated category maps

The out of bounds error paints a similar picture with traditional cGAN suffering from the inability to generate objects within the predefined levels space consistently and this greatly vary among distinct repetitions of the training cycle as the dataset is shuffled differently at each instance. In an attempt to rectify, the network overcorrects which leads to the sensitivity that is noticed from before. There is also other issues that are encountered such as its ability to allot game object type less evenly among the categories by heavily favouring certain categories during the generation process. This is clearly visible in Figure 5.7 where the favoured category keeps varying throughout the multiple independent repetitions of the training algorithm.

5.2.2. Proportions Analysis

By separating the problems, the hybrid architecture with a modified cGAN is shown to better learn the distribution of the various object categories while maintaining appropriate object density when compared to the traditional WGAN-GP as well as the traditional cGAN. Figure 5.7 shows the percentage of objects belonging to each category from sampled levels. While the overall proportions are a major concern in determining the quality of gameplay, there is also a need to moderate the intensity created by the individual object based on the type within the category. As such boss level monsters or the BFG 3000 require a much more appropriate environment to appreciate its presence and must be reflected in the number present in each level. In such cases, the modified cGAN adheres to these principles and sparingly generates objects of higher significance which can be observed in Appendix A that provides average count of each object for every category within the levels.

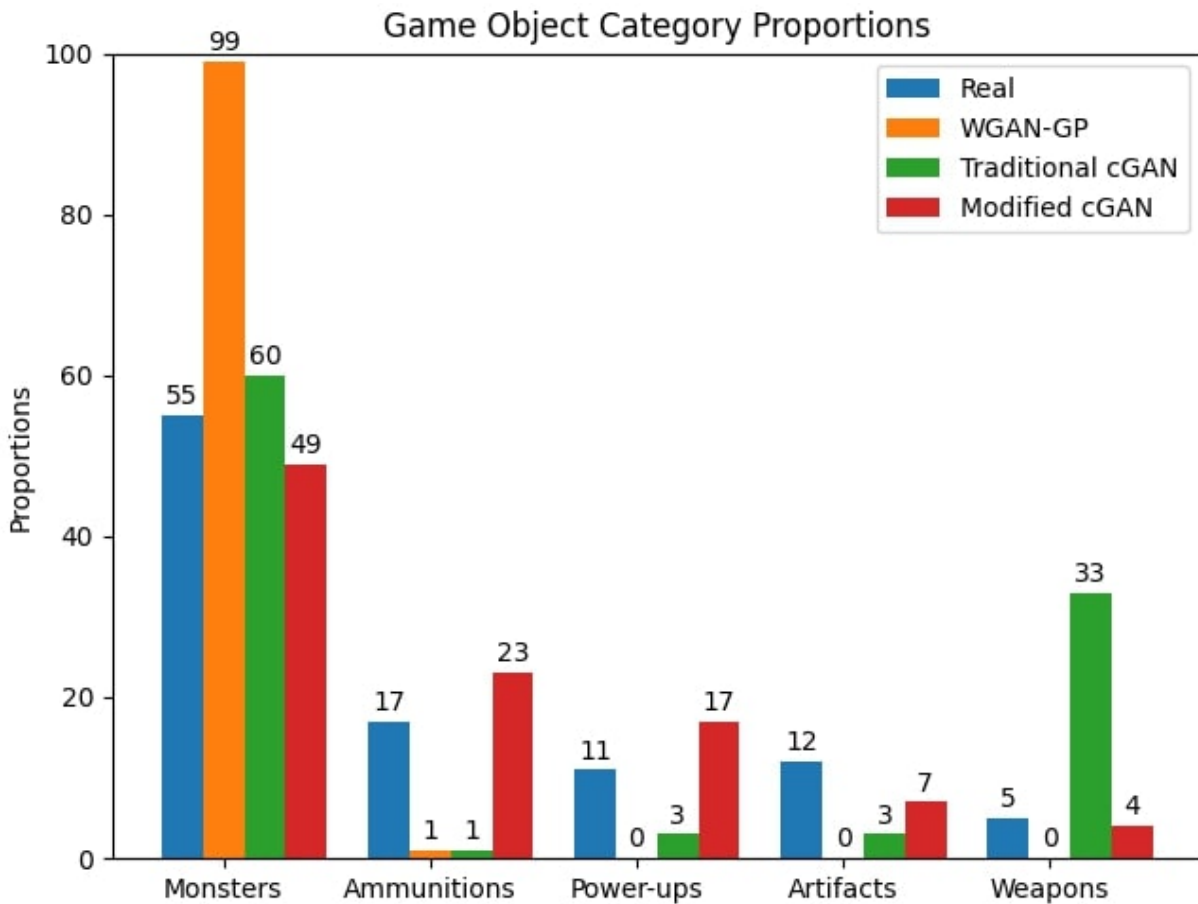


Figure 5.7: Average proportions of each category for sampled levels

5.2.3. Spatial Homogeneity

The Ripley-K results are interpreted with larger values being the outcome of distributions that have greater deviation to completely spatial random distributions or are not spread homogenously across the given space. Table 5.1 records the average spread of the objects with the results closest to the real distribution being that of the Modified cGAN architecture. It is measured at various radii with the histograms of all of the networks at each radius provided in Appendix B. This is done as points have a greater tendency to cluster or spread when focusing on larger or smaller portions than compared to the actual size of the normalized maps [32].

Radius	Real Distribution	WGAN-GP	Traditional cGAN	Modified cGAN
0.5	0.358682	1.044825	0.367623	0.309959
0.6	0.354013	1.125575	0.353378	0.317495
0.7	0.349421	1.137372	0.327674	0.308026
0.8	0.337955	1.098033	0.277786	0.290994
0.9	0.314244	1.014245	0.217251	0.262595
1.0	0.276590	0.877306	0.146254	0.218715

Table 5.1: Average Ripley K value taken at multiple radii

The traditional WGAN-GP network has a greater average Ripley-K result at each of the radii due to overly clumping its game objects near the walls its level structure as well as enlarging the sample space due to the presence of small objects. This is drastically reduced through post processing the topological feature maps as well as using either cGAN networks as it provides a more homogeneous distribution which is in line with that shown by the results from the real feature maps and is why they both lead to a similar result.

5.3. Sample Generation

5.3.1. Traditional WGAN-GP Model




Figure 5.8: Floor map, Height map, Wall map and Things map (columns) generated using the traditional WGAN-GP architecture

As seen in Figure 5.8, the traditional WGAN-GP network is capable of interpreting relations between the wall and floor layouts as well as providing the appropriate height of the various sections but is unable to capture the distribution of the various game objects across the level due to the increased complexity of the problem. As such there occurs a disparity in the distribution of game objects among the categories by mostly generating objects for the first category of the sequenced dictionary id with emphasis towards lower values shown in Figure 5.9 and Figure 5.10. This is better in comparison to training it with the expanded feature maps as it usually generated zero matrices for the categorical features.




Figure 5.9: Floor, Monsters, Ammunitions, Powerups, Artifacts and Weapons Map (Column pairs) generated using the traditional WGAN-GP network (1 of 2)




Figure 5.10: Floor, Monsters, Ammunitions, Powerups, Artifacts and Weapons Map (Column pairs) generated using the traditional WGAN-GP network (2 of 2)

5.3.2. Hybrid Model with Traditional cGAN




Figure 5.11: Floor, Monsters, Ammunitions, Powerups, Artifacts and Weapons Map (Column pairs) generated using the hybrid architecture with traditional cGAN network (1 of 2)




Figure 5.12: Floor, Monsters, Ammunitions, Powerups, Artifacts and Weapons Map (Column pairs) generated using the hybrid architecture with traditional cGAN network (2 of 2)

5.3.3. Hybrid Model with Modified cGAN




Figure 5.13: Floor map, Height map, Wall map and Things map (rows) generated using the hybrid architecture with modified cGAN network

The hybrid network works around this by separating it into 2 distinct subproblems, allowing the WGAN-GP to focus primarily on the level structure while the cGAN network maps the object positions from the generated level structure. This also allows the WGAN-GP to focus more resources towards learning the topological distribution which helps in the development of better level layouts to a certain extent. This does not ensure that its structure is always better given a random seed as the samples that have been use are handpicked to showcase the best results of all the trained architectures. As both hybrid models use the same WGAN-GP generator, the topological maps are the same with the only difference lying in its generated object maps.




Figure 5.14: Floor, Monsters, Ammunitions, Powerups, Artifacts and Weapons Map (Column pairs) generated using the hybrid architecture with modified cGAN network (1 of 2)




Figure 5.15: Floor, Monsters, Ammunitions, Powerups, Artifacts and Weapons Map (Column pairs) generated using the hybrid architecture with modified cGAN network (2 of 2)

5.3.4. Condensed Category Map Comparison

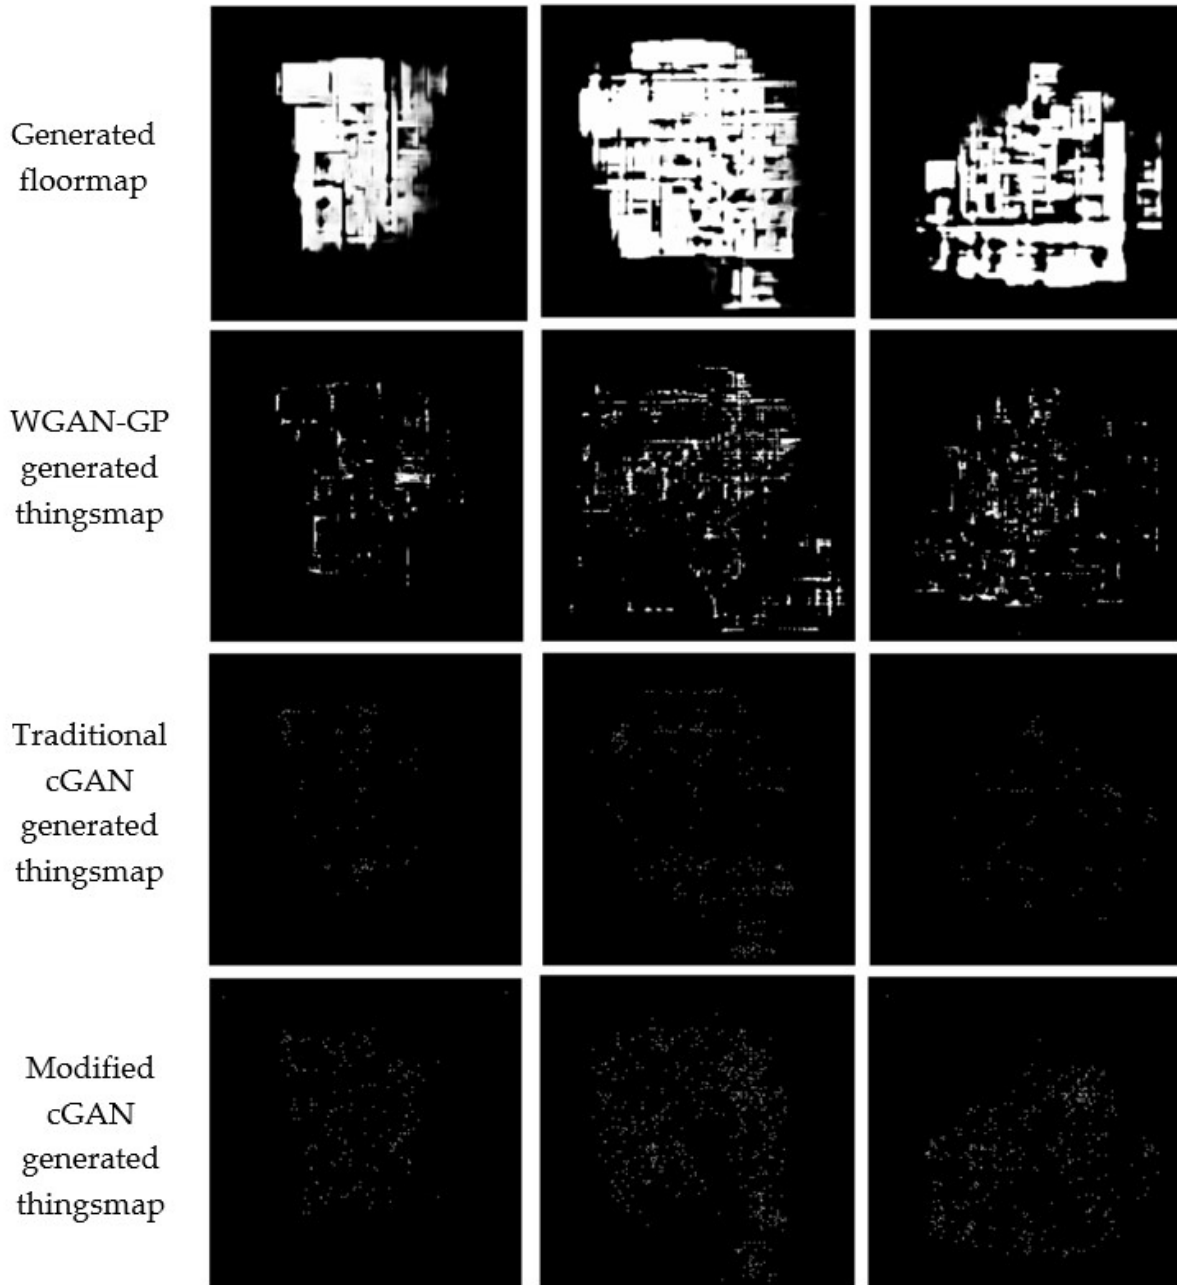


Figure 5.16: Generated things maps of the various models using levels generated by the WGAN-GP

6 | Conclusion

The proposed system should serve as an early prototype for a more holistic format of procedural generating 3D levels using generative adversarial networks. The hybrid system has shown its prowess to suitably provide objects within a given level space and has solved any issues that could hinder gameplay from the previous design. It has proven its ability to provide the minimum subset of features needed to capture all the fundamental aspects of a level that can be exercised in isometric game engines by taking advantage of the simplicity of level structure in earlier games. This can be replicated in most games with a 2D map system to improve their replayability through the generation of multiple level layouts as well as multiple game object maps for each layout occasionally.

The models in this experiment are restrained to using the same dataset for the sake of maintaining consistency across all of the trained architectures. The system is open to improvement by modifying the dataset so that each model within the hybrid architecture can be fully utilized to its utmost potential. This can be done simply by altering the dataset that is used to train each of the models by tailoring it towards their strengths by means of reducing the image size and dropping the padding entirely for the WGAN-GP subsystem while making use of this preprocessing for the cGAN with larger maps. This model can also incorporate other categories present within DOOM object types such as the obstacles and decorations.

While there are certain limitations and room for improvement, it stands as a better alternative to replaying levels in games, especially when they rely on ‘grinding’ the same maps to slow the pace at which the player progresses through its content to avoid them from finishing the game too quickly as it might indirectly affect their brand image. There is also the added benefit of reducing excessive dependence on content developers in designing every aspect of all the level and rather provide intricate details to the ones that are cinematic in nature. This would maximize content in both quality and quantity by relying on such systems to provide generic levels as buffers between impactful set pieces and allow for an all around better pacing within games.

6.1. Limitations

There are a few constraints regarding its applicability to other games as it requires game data to be stored in a format where level structures can be easily decoded and parsed into image maps manually or through level builders with detailed specifications or guides. This system is also dependent on the game having a large collection of level instances created either by the studio or through effective tools for the community to build, access and share. Games are limited to simple terrains as it is unable to represent multiple floors or stacked sections which are required to construct native 3D maps. It does not allow for different noisy images to consistently generate significantly different category maps as there are usually very slight differences in their placements. There is also an issue in implementing doors and keys which need to be addressed as well as the use of teleporters to connect detached segments within the level layout by including the respective triggers and landings in the generator before they can be used to incorporate concepts of puzzles within gameplay.

6.2. Future Works

This system can be expanded by implementing additional models to separately learn the texture distribution of the level with the texture maps implemented through encoding pixels with texture id from a graphic dictionary and generated by means of a conditional network as used here. Conversely, by modifying the architecture to represent the textures map at lower resolutions, actual pixel colors can be encoded in the images to generate new graphics that are then extracted from the maps and encoded into their respective formats. Although this would require image translation with super resolution models [41] that can generate the aesthetical features together as the process would be dramatically upscaled for representing the complete texture map at an acceptable resolution.

By adding relations between pairs of objects such as teleporters and landings or locked doors and keys that can be introduced through a separate feature map representation, it can be used to improve the versatility of the generator to be used in other genres as long as steps are taken to ensure that the levels generated meet the standard requirements such as solvability through additions made in the shortest path algorithms when considering these objects. The system can also be modified with different sets of objects and used to train agents within simulated environments by learning a possible set of actions or behaviors given certain objects present in them.

Bibliography

- [1] M. Arjovsky, S. Chintala, and L. Bottou. Wasserstein gan. *ICML'17: Proceedings of the 34th International Conference on Machine Learning*, 70:214–223, 08 2017.
- [2] C. Beckham and C. Pal. A step towards procedural terrain generation with gans. *34th International Conference on Machine Learning(ICML)*, 08 2017.
- [3] D. Braben and I. Bell. Elite, 1984. game.
- [4] Y. Chen, I. Keivanloo, and C. Roy. Near-miss software clones in open source games: An empirical study. *Canadian Conference on Electrical and Computer Engineering*, pages 1–7, 05 2014. doi: 10.1109/CCECE.2014.6901018.
- [5] S. Earle. Using fractal neural networks to play simcity 1 and conway’s game of life at variable scales, 01 2020. URL <https://arxiv.org/abs/2002.03896>.
- [6] M. Fell. The unofficial doom specs, 1994. URL <https://www.gamers.org/dhs/helpdocs/dmsp1666.htmlc>.
- [7] T. Frattesi, D. Griesbach, J. Leith, and T. Shaffer. Replayability of video games. Master’s thesis, Worcester Polytechnic Institute, Worcestershire, England, 05 2011.
- [8] H. Games. No man’s sky, 2016. game.
- [9] E. Giacomello, P. L. Lanzi, and D. Loiacono. Doom level generation using generative adversarial networks. pages 316–323, 2018. doi: 10.1109/GEM.2018.8516539.
- [10] I. Goodfellow, J. Pouget-Abadie, M. Mirza, B. Xu, D. Warde-Farley, S. Ozair, A. Courville, and Y. Bengio. Generative adversarial networks. *Advances in Neural Information Processing Systems*, 3, 06 2014. doi: 10.1145/3422622.
- [11] I. Gulrajani, F. Ahmed, M. Arjovsky, V. Dumoulin, and A. Courville. Improved training of wasserstein gans. *NIPS’17: Proceedings of the 31st International Conference on Neural Information Processing Systems*, page 5769–5779, 12 2017.
- [12] M. Hendrikx, S. Meijer, J. Velden, and A. Iosup. Procedural content generation for

- games: A survey. *ACM Transactions on Multimedia Computing, Communications, and Applications (TOMCCAP)*, 9, 02 2013. doi: 10.1145/2422956.2422957.
- [13] S. Hochreiter and J. Schmidhuber. Long short-term memory. *Neural computation*, 9:1735–80, 12 1997. doi: 10.1162/neco.1997.9.8.1735.
- [14] id Software. Doom, 1993. game.
- [15] P. Isola, J.-Y. Zhu, T. Zhou, and A. Efros. Image-to-image translation with conditional adversarial networks. pages 5967–5976, 07 2017. doi: 10.1109/CVPR.2017.632. URL <https://doi.ieeecomputersociety.org/10.1109/CVPR.2017.632>.
- [16] S. Karakovskiy and J. Togelius. The mario ai benchmark and competitions. *Computational Intelligence and AI in Games, IEEE Transactions on*, 4:55–67, 03 2012. doi: 10.1109/TCIAIG.2012.2188528.
- [17] D. Karavolos, A. Liapis, and G. Yannakakis. A multifaceted surrogate model for search-based procedural content generation. *IEEE Transactions on Games*, 13(1): 11–22, 2021. doi: 10.1109/TG.2019.2931044.
- [18] T. Karras, S. Laine, and T. Aila. A style-based generator architecture for generative adversarial networks. *CoRR*, abs/1812.04948:4396–4405, 06 2019. doi: 10.1109/CVPR.2019.00453. URL <http://arxiv.org/abs/1812.04948>.
- [19] A. Khalifa, P. Bontrager, S. Earle, and J. Togelius. Pcgrl: Procedural content generation via reinforcement learning. *CoRR*, abs/2001.09212, 01 2020. URL <https://arxiv.org/abs/2001.09212>.
- [20] G. Lample and D. Chaplot. Playing fps games with deep reinforcement learning. *CoRR*, abs/1609.05521, 09 2016. URL <http://arxiv.org/abs/1609.05521>.
- [21] Z. Li, S. Zhu, B. Shao, X. Zeng, T. Wang, and T.-Y. Liu. Dsn-ddi: an accurate and generalized framework for drug–drug interaction prediction by dual-view representation learning. *Briefings in Bioinformatics*, 24, 01 2023. doi: 10.1093/bib/bbac597.
- [22] J. Liu, S. Snodgrass, A. Khalifa, S. Risi, G. Yannakakis, and J. Togelius. Deep learning for procedural content generation. *Neural Computing and Applications*, 33(1):19–37, 10 2020. doi: 10.1007/s00521-020-05383-8. URL <https://doi.org/10.1007/2Fs00521-020-05383-8>.
- [23] W. McCulloch and W. Pitts. A logical calculus of the ideas immanent in nervous activity. *Bulletin of Mothematicnl Biology*, 52(1/2):99–115, 1990. URL <http://arxiv.org/abs/1706.03762>.

- [24] MicroProse. Civilization ii, 1996. game.
- [25] K. Park, B. W. Mott, W. Min, K. E. Boyer, E. N. Wiebe, and J. C. Lester. Generating educational game levels with multistep deep convolutional generative adversarial networks. pages 1–8, 2019. doi: 10.1109/CIG.2019.8848085.
- [26] A. Radford, L. Metz, and S. Chintala. Unsupervised representation learning with deep convolutional generative adversarial networks. 11 2016. URL <http://arxiv.org/abs/1511.06434>.
- [27] B. D. Ripley. The second-order analysis of stationary point processes. *Journal of Applied Probability*, 13:255–266, 06 1976. doi: 10.2307/3212829.
- [28] R. Rombach, A. Blattmann, D. Lorenz, P. Esser, and B. Ommer. High-resolution image synthesis with latent diffusion models. pages 10674–10685, 06 2022. doi: 10.1109/CVPR52688.2022.01042. URL <https://doi.ieeecomputersociety.org/10.1109/CVPR52688.2022.01042>.
- [29] O. Ronneberger, P. Fischer, and T. Brox. U-net: Convolutional networks for biomedical image segmentation. *LNCS*, 9351:234–241, 10 2015. doi: 10.1007/978-3-319-24574-4_28.
- [30] M. Rousseau. Zennode, 2004. URL <https://www.mrousseau.org/programs/ZenNode/>. node,blockmap,reject builder.
- [31] A. Sauer, K. Chitta, J. Muller, and A. Geiger. Projected gans converge faster. *CoRR*, abs/2111.01007, 2021. URL <https://arxiv.org/abs/2111.01007>.
- [32] J. Sparring, R. Waagepetersen, and S. Sommer. Generalizations of ripley’s k-function with application to space curves. *26th International Conference on Information Processing in Medical Imaging (IPMI)*, 11492:731–742, 12 2018.
- [33] A. Summerville, M. Guzdial, M. Mateas, and M. Riedl. Learning player tailored content from observation: Platformer level generation from video traces using lstms. 12:107–113, 06 2021. doi: 10.1609/aiide.v12i2.12895. URL <https://ojs.aaai.org/index.php/AIIDE/article/view/12895>.
- [34] D. Talwar, S. Guruswamy, N. Ravipati, and M. Eirinaki. Evaluating validity of synthetic data in perception tasks for autonomous vehicles. pages 73–80, 08 2020. doi: 10.1109/AITEST49225.2020.00018.
- [35] J. Togelius, G. N. Yannakakis, K. O. Stanley, and C. Browne. Search-based procedural content generation: A taxonomy and survey. *IEEE Transactions on Compu-*

- tational Intelligence and AI in Games*, 3(3):172–186, 2011. doi: 10.1109/TCIAIG.2011.2148116.
- [36] M. Toy and G. Wichman. Rogue, 1980. game.
- [37] Y. Tsujino and R. Yamanishi. Dance dance gradation: A generation of fine-tuned dance charts. *17th IFIP TC 14 International Conference, WCC 2018*, pages 175–187, 08 2018. doi: 10.1007/978-3-319-99426-0_15.
- [38] T. Tutenel, R. M. Smelik, R. Lopes, K. J. de Kraker, and R. Bidarra. Generating consistent buildings: A semantic approach for integrating procedural techniques. *IEEE Transactions on Computational Intelligence and AI in Games*, 3(3):274–288, 2011. doi: 10.1109/TCIAIG.2011.2162842.
- [39] A. Vaswani, N. Shazeer, N. Parmar, J. Uszkoreit, L. Jones, A. Gomez, L. Kaiser, and I. Polosukhin. Attention is all you need. *CoRR*, abs/1706.03762, 06 2017. URL <http://arxiv.org/abs/1706.03762>.
- [40] C.-Y. Wang, I.-H. Yeh, and H.-y. Liao. You only learn one representation: Unified network for multiple tasks. *CoRR*, abs/2105.04206, 05 2021. URL <https://arxiv.org/abs/2105.04206>.
- [41] X. Wang, J. Yi, J. Guo, Y. Song, J. Lyu, J. Xu, W. Yan, J. Zhao, Q. Cai, and H. Min. A review of image super-resolution approaches based on deep learning and applications in remote sensing. *Remote Sensing*, 14(21), 2022. ISSN 2072-4292. doi: 10.3390/rs14215423. URL <https://www.mdpi.com/2072-4292/14/21/5423>.
- [42] L. Wilschut, A. Laudisoit, N. Hughes, E. Addink, S. De Jong, H. Heesterbeek, J. Reijnders, S. Eagle, V. Dubyanskiy, and M. Begon. Spatial distribution patterns of plague hosts: Point pattern analysis of the burrows of great gerbils in kazakhstan. *Journal of Biogeography*, 42, 05 2015. doi: 10.1111/jbi.12534.
- [43] L. Yang, Z. Zhang, Y. Song, S. Hong, R. Xu, Y. Zhao, W. Zhang, B. Cui, and M.-H. Yang. Diffusion models: A comprehensive survey of methods and applications. 09 2022. doi: 10.48550/arXiv.2209.00796.
- [44] A. Zeyer, P. Bahar, K. Irie, R. Schluter, and H. Ney. A comparison of transformer and lstm encoder decoder models for asr. *IEEE Automatic Speech Recognition and Understanding*, pages 8–15, 12 2019. doi: 10.1109/ASRU46091.2019.9004025.
- [45] J.-Y. Zhu, T. Park, P. Isola, and A. A. Efros. Unpaired image-to-image translation using cycle-consistent adversarial networks. pages 2242–2251, 2017. doi: 10.1109/ICCV.2017.244.

A | Appendix A

Ripley K Histograms




Figure A.1: Ripley K histogram with radius 0.5




Figure A.2: Ripley K histogram with radius 0.6




Figure A.3: Ripley K histogram with radius 0.7




Figure A.4: Ripley K histogram with radius 0.8




Figure A.5: Ripley K histogram with radius 0.9




Figure A.6: Ripley K histogram with radius 1.0

B | Appendix B

Game Objects Count by Type




Figure B.1: Average monsters count by type in a level




Figure B.2: Average ammunitions count by type in a level




Figure B.3: Average powerups count by type in a level




Figure B.4: Average artifacts count by type in a level




Figure B.5: Average weapons count by type in a level

List of Figures

2.1	Generated numbers and faces from the work of Goodfellow et al.	8
2.2	Generated bedrooms from the work of Radford et al.	9
2.3	Generated Google Maps images from the work of Isola et al.	12
3.1	Cloned lines of code for C-based from the work of Chen et al.	13
3.2	Sample DOOM PWAD visualized using DOOM Builder	17
3.3	Feature maps of a sample level	19
3.4	Conditional WGAN-GP architecture from the work of Giacomello et al. . .	20
3.5	Resized feature maps comparison	22
3.6	Area distribution of the filtered dataset	23
3.7	Updated feature maps of a sample level	24
4.1	Hybrid GAN architecture	28
4.2	Topological generator of the hybrid architecture	29
4.3	Functional generator of the hybrid architecture	31
4.4	Data flow diagram	32
5.1	Discriminator Loss of the traditional cGAN model	41
5.2	Discriminator Loss of the modified cGAN model	42
5.3	Mean entropy of the generated category maps	43
5.4	Mean object count per unit area of the generated category maps	43
5.5	Mean encoding error of the generated category maps	44
5.6	Mean out of bounds error of the generated category maps	45
5.7	Average proportions of each category for sampled levels	46
5.8	Samples generated using the traditional WGAN-GP architecture	48
5.9	Samples generated using the traditional WGAN-GP architecture (1 of 2) .	49
5.10	Samples generated using the traditional WGAN-GP architecture (2 of 2) .	50
5.11	Samples generated using the traditional cGAN network (1 of 2)	51
5.12	Samples generated using the traditional cGAN network (2 of 2)	52
5.13	Samples generated using the modified cGAN network	53
5.14	Samples generated using the modified cGAN network (1 of 2)	54

5.15	Samples generated using the modified cGAN network (2 of 2)	55
5.16	Generated things maps comparison	56
A.1	Ripley K histogram with radius 0.5	63
A.2	Ripley K histogram with radius 0.6	64
A.3	Ripley K histogram with radius 0.7	64
A.4	Ripley K histogram with radius 0.8	65
A.5	Ripley K histogram with radius 0.9	65
A.6	Ripley K histogram with radius 1.0	66
B.1	Average monsters count by type in a level	67
B.2	Average ammunitions count by type in a level	68
B.3	Average powerups count by type in a level	68
B.4	Average artifacts count by type in a level	69
B.5	Average weapons count by type in a level	69

List of Tables

3.1	Sections of a WAD as seen in the work of Giacomello et al.	14
3.2	Unique mandatory lumps in a WAD	15
3.3	Derivable mandatory lumps in a WAD	16
3.4	List of relevant game objects with category ids (1 of 2)	25
3.5	List of relevant game objects with category ids (2 of 2)	26
4.1	Training operation of WGAN-GP models	34
4.2	Training operation of cGAN models	35
4.3	Trained architectures used for the comparative study	37
5.1	Average Ripley K value taken at multiple radii	47

Acknowledgements

I would like to thank my supervisor, Daniele Loiacono for the instructions he provided me that allowed me to write this thesis. I thank my family and friends, especially my parents who supported me both financially and emotionally through this prolonged endeavor as well as provided me with their consistent encouragement which kept me motivated and stay on track towards completing this thesis. I express my gratitude towards Ali Emre Soleleyuk that sat through my venting during arising complications which let me have the breakthroughs that helped me progress towards its completion and Amrit A. with whom I was able to relax and redirect my focus towards actions that have been more productive.

